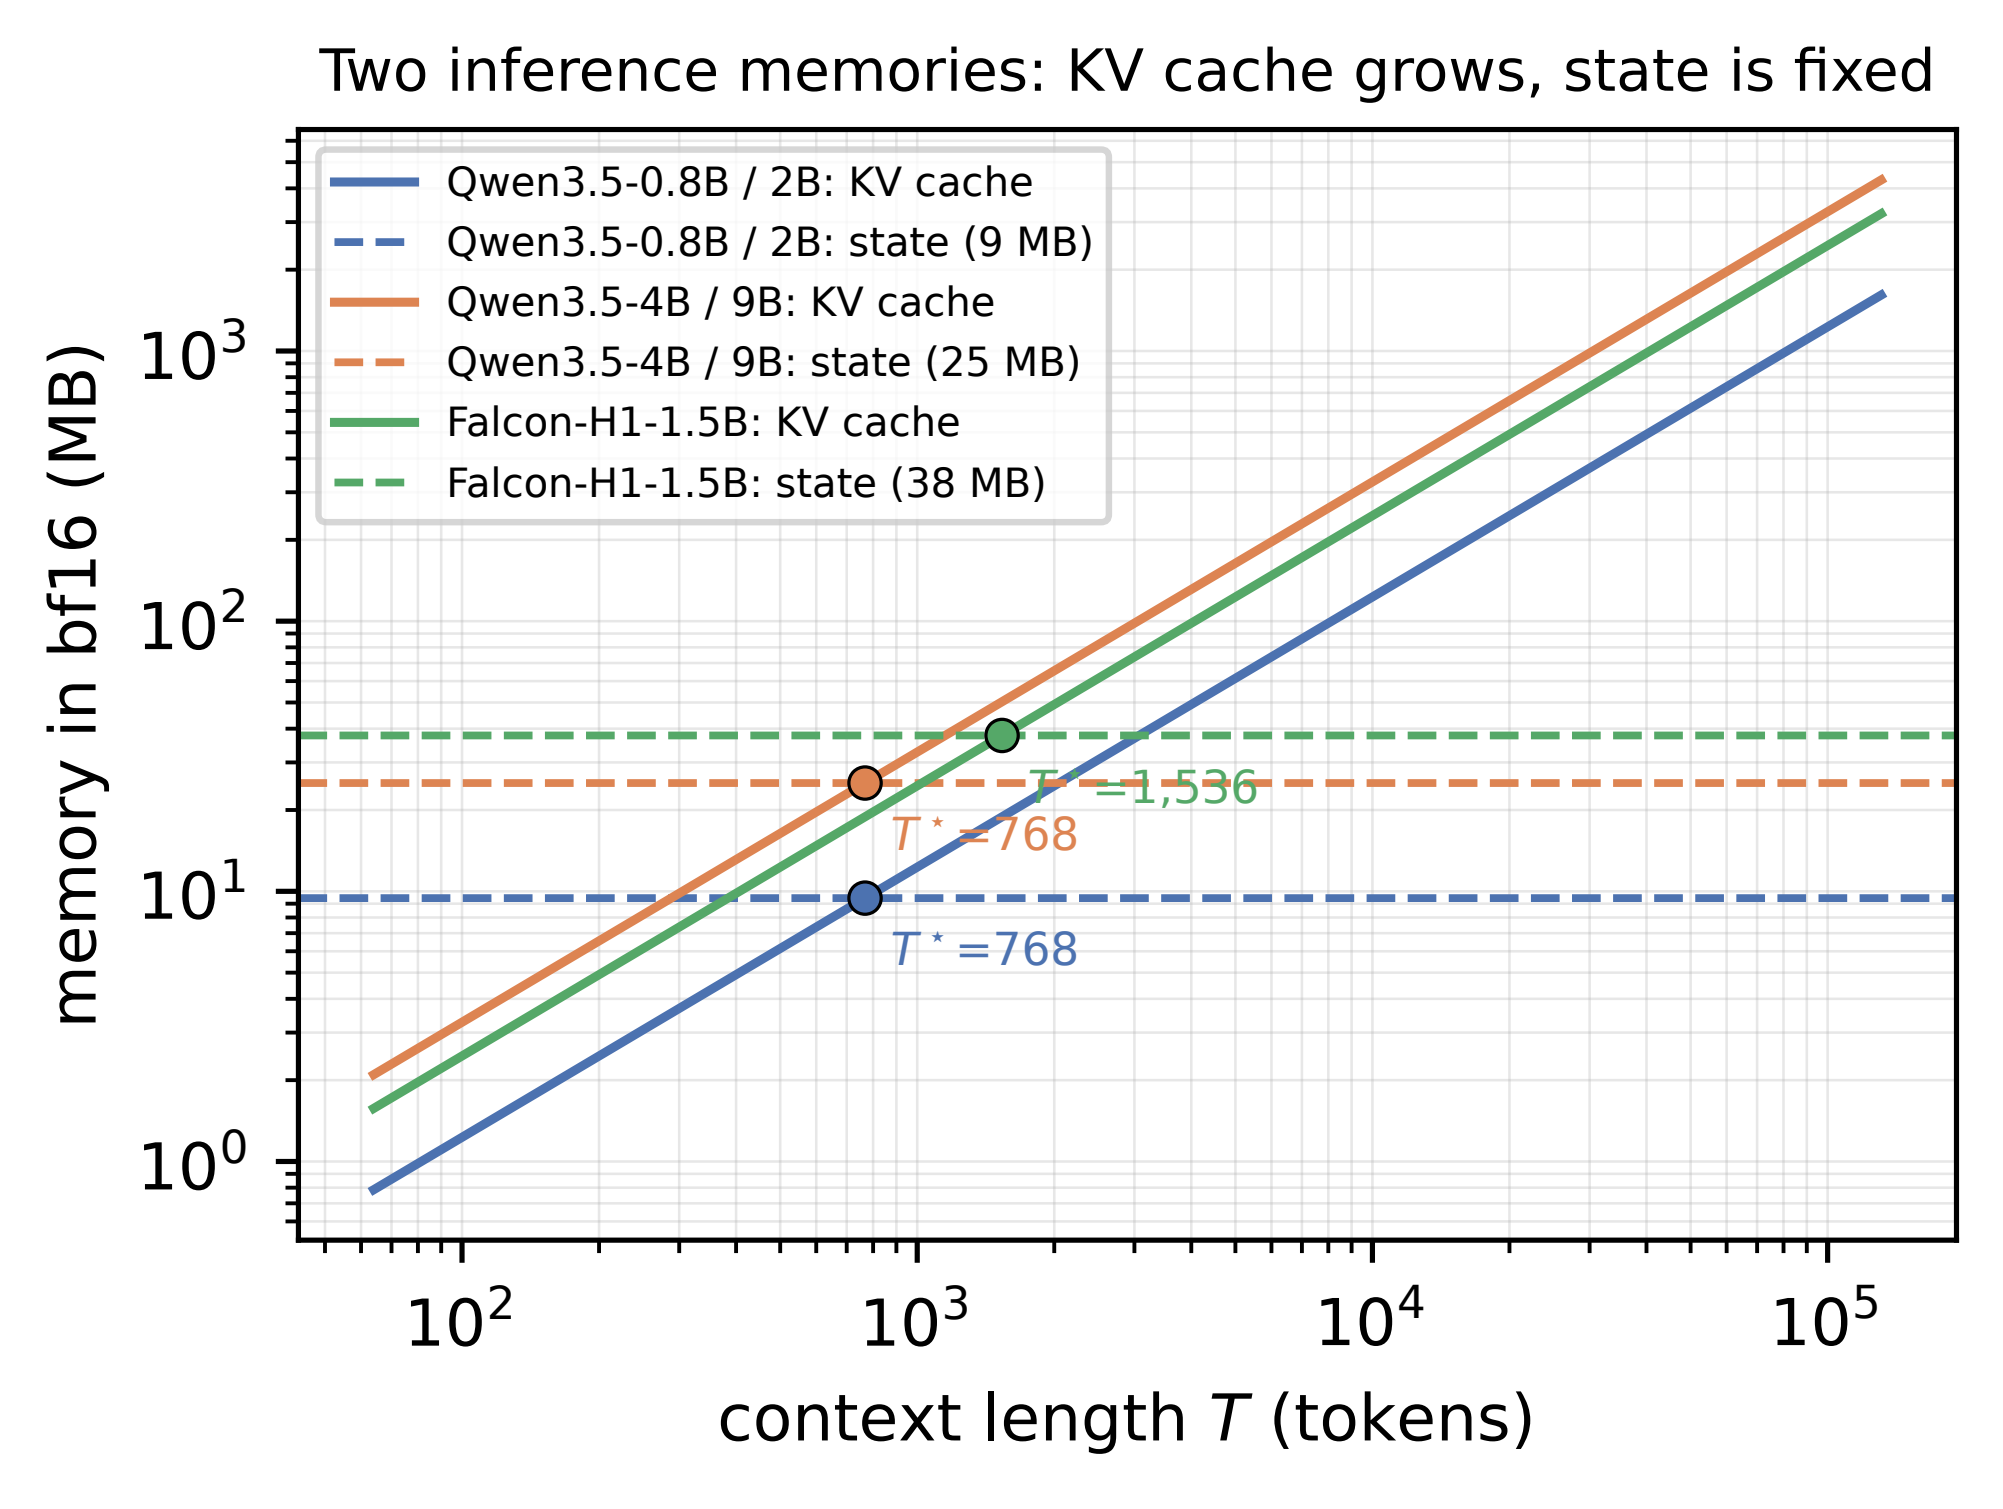
Two inference memories: KV cache grows, state is fixed
Qwen3.5-0.8B / 2B: KV cache
Qwen3.5-0.8B / 2B: state (9 MB)
Qwen3.5-4B / 9B: KV cache
Qwen3.5-4B / 9B: state (25 MB)
Falcon-H1-1.5B: KV cache
Falcon-H1-1.5B: state (38 MB)
$T^\star$ =1,536
$T^\star$ =768
$T^\star$ =768
memory in bf16 (MB)
10^0
10^1
10^2
10^3
10^2
10^3
10^4
10^5
context length T (tokens)